

Attribute information For bank dataset

Input variables:

- bank client data:
- 1 - age (numeric)
 - 2 - job : type of job (categorical: "admin.", "unknown", "unemployed", "management", "housemaid", "entrepreneur", "student", "blue-collar", "self-employed", "retired", "technician", "services")
 - 3 - marital : marital status (categorical: "married", "divorced", "single"; note: "divorced" means divorced or widowed)
 - 4 - education (categorical: "unknown", "secondary", "primary", "tertiary")
 - 5 - default: has credit in default? (binary: "yes", "no")
 - 6 - balance: average yearly balance, in euros (numeric)
 - 7 - housing: has housing loan? (binary: "yes", "no")
 - 8 - loan: has personal loan? (binary: "yes", "no")

related with the last contact of the current campaign:

- 9 - contact: contact communication type (categorical: "unknown", "telephone", "cellular")
- 10 - day: last contact day of the month (numeric)
- 11 - month: last contact month of year (categorical: "jan", "feb", "mar", ..., "nov", "dec")
- 12 - duration: last contact duration, in seconds (numeric)

other attributes:

- 13 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- 14 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)
- 15 - previous: number of contacts performed before this campaign and for this client (numeric)
- 16 - poutcome: outcome of the previous marketing campaign (categorical: "unknown", "other", "failure", "success")

Output variable (desired target):

- 17 - y - has the client subscribed a term deposit? (binary: "yes", "no")

8. Missing Attribute Values: None

```
In [33]: # Importing the Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: bank_data=pd.read_csv('bank-full.csv', sep = ';')
bank_data.head()
```

Out[3]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previou
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	

EDA

```
In [4]: bank_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays      45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

```
In [5]: bank_data.dtypes
```

```
Out[5]: age          int64
job            object
marital        object
education      object
default        object
balance        int64
housing        object
loan           object
contact        object
day            int64
month          object
duration       int64
campaign       int64
pdays         int64
previous       int64
poutcome      object
y              object
dtype: object
```

```
In [6]: bank_data.describe()
```

```
Out[6]:
```

	age	balance	day	duration	campaign	pdays	previous
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.936210	1362.272058	15.806419	258.163080	2.763841	40.197828	0.580323
std	10.618762	3044.765829	8.322476	257.527812	3.098021	100.128746	2.303441
min	18.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000	0.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000	-1.000000	0.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000	-1.000000	0.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000	0.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000	871.000000	275.000000

```
In [7]: bank_data.isna().sum()
```

```
Out[7]: age          0
job            0
marital        0
education      0
default        0
balance        0
housing        0
loan           0
contact        0
day            0
month          0
duration       0
campaign       0
pdays         0
previous       0
poutcome      0
y              0
dtype: int64
```

In [8]: ▶ bank_data.nunique()

Out[8]: age 77
job 12
marital 3
education 4
default 2
balance 7168
housing 2
loan 2
contact 3
day 31
month 12
duration 1573
campaign 48
pdays 559
previous 41
poutcome 4
y 2
dtype: int64

Data Preprocessing & Visualization

In [9]: ▶ *# Renaming target variable 'y' to 'Deposit' and moving it to the first position*
dep = bank_data['y']
#Drop the deposit column
bank_data.drop(labels=['y'], axis=1,inplace = True)
bank_data.insert(0, 'Deposit', dep)
bank_data.head()

Out[9]:

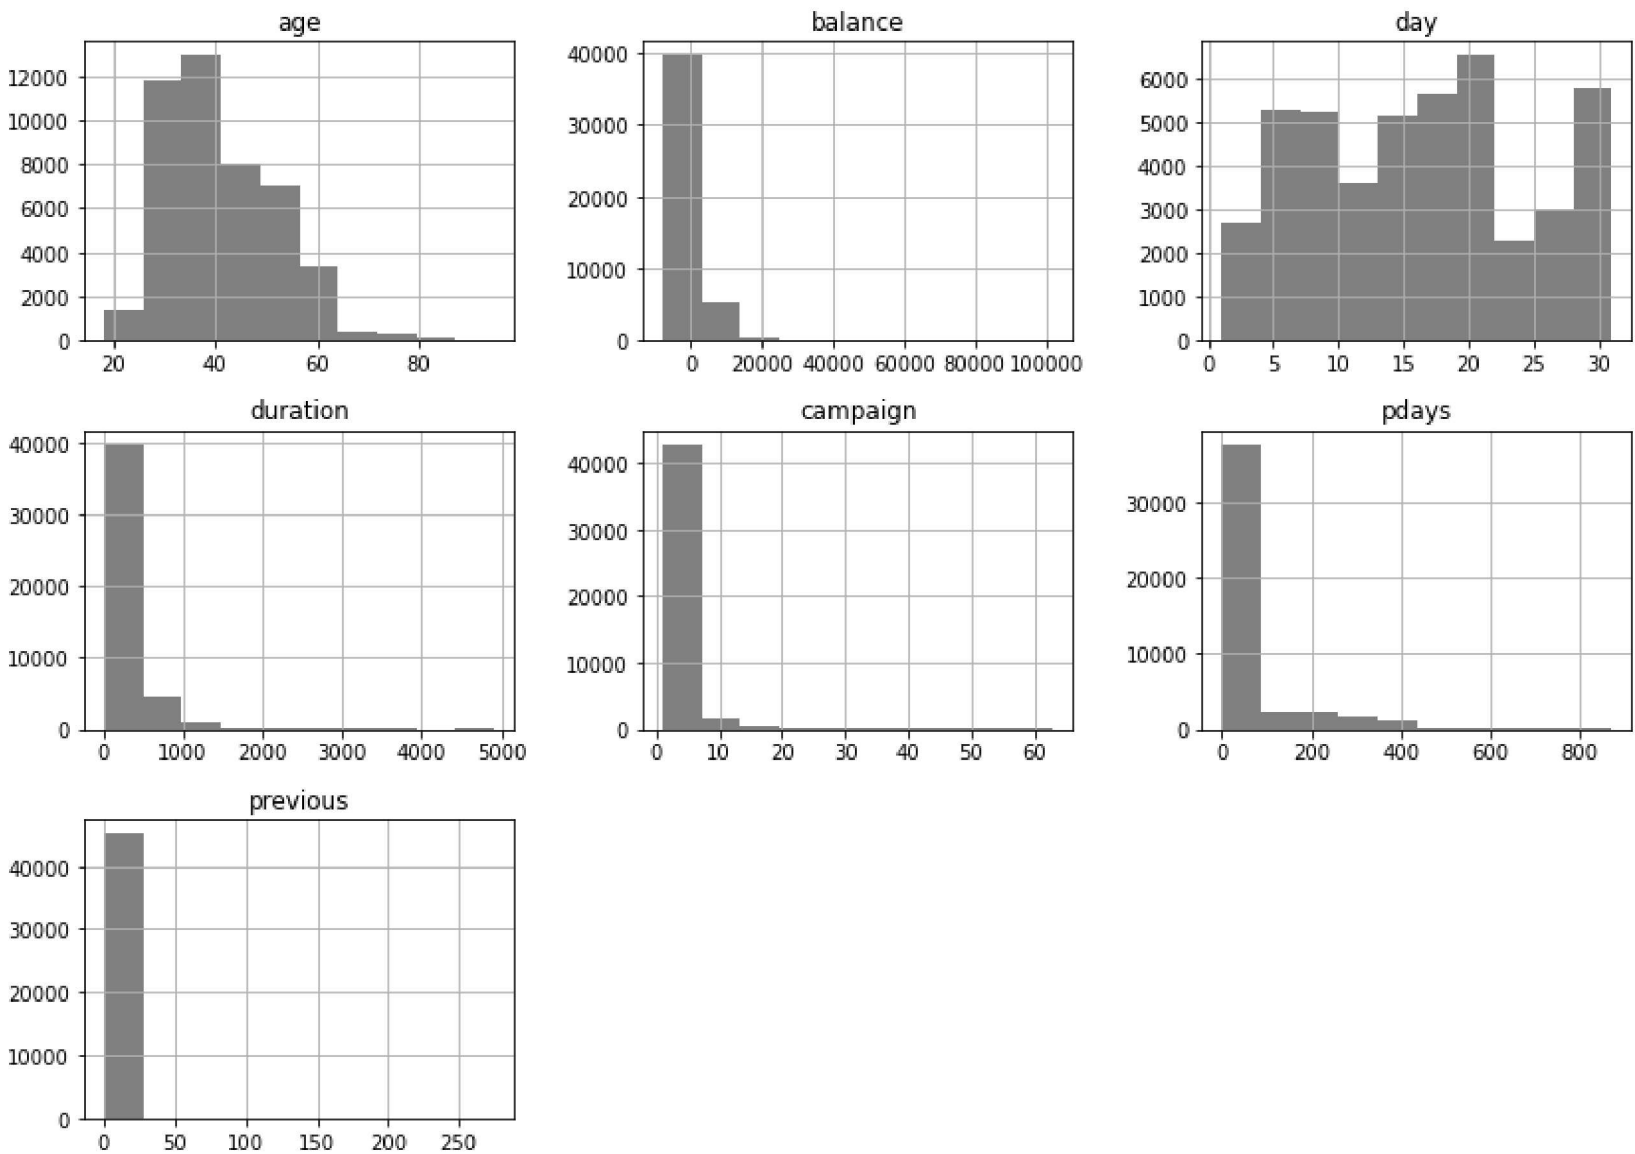
	Deposit	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays
0	no	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1
1	no	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1
2	no	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1
3	no	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1
4	no	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1

In [10]: ▶ bank_data['Deposit'].value_counts()

Out[10]: no 39922
yes 5289
Name: Deposit, dtype: int64

In [12]: ▶ *# Plotting numeric data distribution.*

bank_data.hist(bins=10, figsize=(14,10), color='grey')
plt.show()



```
In [14]: ▶ # One-Hot Encoding of categrical variables
data=pd.get_dummies(bank_data,columns=['job','marital','education','contact','poutcome'])
pd.set_option("display.max.columns", None)
data.head()
```

Out[14]:

	Deposit	age	default	balance	housing	loan	day	month	duration	campaign	pdays	previous	job_admin.	job_blue-collar	job_entr
0	no	58	no	2143	yes	no	5	may	261	1	-1	0	0	0	
1	no	44	no	29	yes	no	5	may	151	1	-1	0	0	0	
2	no	33	no	2	yes	yes	5	may	76	1	-1	0	0	0	
3	no	47	no	1506	yes	no	5	may	92	1	-1	0	0	1	
4	no	33	no	1	no	no	5	may	198	1	-1	0	0	0	

```
In [15]: ▶ # Convert the columns that contain a Yes or No. (Binary Columns)
def convert_to_int(data, new_column, target_column):
    data[new_column] = data[target_column].apply(lambda x: 0 if x == 'no' else 1)
    return data[new_column].value_counts()
```

```
In [16]: ▶ convert_to_int(data, "deposit_int", "Deposit") #Create a deposit int
convert_to_int(data, "housing_int", "housing") # Create housingint column
convert_to_int(data, "loan_int", "loan") #Create a Loan_int column
convert_to_int(data, "default_int", "default") #Create a default_int column
```

Out[16]: 0 44396
1 815
Name: default_int, dtype: int64

```
In [17]: ▶ # Drop the binary columns and Leave the same column in the form of integers 0 = No and 1 = Yes
data.drop(['housing', 'loan', 'default'], axis=1, inplace=True)
```

```
In [18]: ▶ # Find and Replace Encoding for month categorical variable
data['month'].value_counts()
```

Out[18]: may 13766
jul 6895
aug 6247
jun 5341
nov 3970
apr 2932
feb 2649
jan 1403
oct 738
sep 579
mar 477
dec 214
Name: month, dtype: int64

```
In [19]: ▶ order={'month':{'jan':1,'feb':2,'mar':3,'apr':4,'may':5,'jun':6,'jul':7,'aug':8,'sep':9,'oct':10,'nov':11,'d
```

```
In [20]: ▶ data=data.replace(order)
data.head()
```

Out[20]:

	Deposit	age	balance	day	month	duration	campaign	pdays	previous	job_admin.	job_blue-collar	job_entrepreneur	job_housemaic
0	no	58	2143	5	5	261	1	-1	0	0	0	0	(
1	no	44	29	5	5	151	1	-1	0	0	0	0	(
2	no	33	2	5	5	76	1	-1	0	0	0	1	(
3	no	47	1506	5	5	92	1	-1	0	0	1	0	(
4	no	33	1	5	5	198	1	-1	0	0	0	0	(

```
In [21]: ▶ data.drop(['Deposit'], axis=1, inplace=True)
# Rename deposit_int column for Deposit and then move it to the first
data = data.rename(columns={"deposit_int": "deposit"})
first = data['deposit']
data.drop(labels=['deposit'], axis=1,inplace = True)
# insert (loc, column, values) --> loc is the same as position in the column.
data.insert(0, 'deposit', first)
data["deposit"].value_counts()
```

Out[21]: 0 39922
1 5289
Name: deposit, dtype: int64

```
In [22]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 38 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   deposit                              45211 non-null  int64
1   age                                  45211 non-null  int64
2   balance                             45211 non-null  int64
3   day                                  45211 non-null  int64
4   month                               45211 non-null  int64
5   duration                             45211 non-null  int64
6   campaign                            45211 non-null  int64
7   pdays                               45211 non-null  int64
8   previous                            45211 non-null  int64
9   job_admin.                           45211 non-null  uint8
10  job_blue-collar                      45211 non-null  uint8
11  job_entrepreneur                     45211 non-null  uint8
12  job_housemaid                        45211 non-null  uint8
13  job_management                       45211 non-null  uint8
14  job_retired                          45211 non-null  uint8
15  job_self-employed                   45211 non-null  uint8
16  job_services                         45211 non-null  uint8
17  job_student                          45211 non-null  uint8
18  job_technician                      45211 non-null  uint8
19  job_unemployed                      45211 non-null  uint8
20  job_unknown                         45211 non-null  uint8
21  marital_divorced                    45211 non-null  uint8
22  marital_married                     45211 non-null  uint8
23  marital_single                      45211 non-null  uint8
24  education_primary                   45211 non-null  uint8
25  education_secondary                 45211 non-null  uint8
26  education_tertiary                  45211 non-null  uint8
27  education_unknown                   45211 non-null  uint8
28  contact_cellular                    45211 non-null  uint8
29  contact_telephone                   45211 non-null  uint8
30  contact_unknown                     45211 non-null  uint8
31  poutcome_failure                    45211 non-null  uint8
32  poutcome_other                      45211 non-null  uint8
33  poutcome_success                    45211 non-null  uint8
34  poutcome_unknown                   45211 non-null  uint8
35  housing_int                         45211 non-null  int64
36  loan_int                            45211 non-null  int64
37  default_int                         45211 non-null  int64
dtypes: int64(12), uint8(26)
memory usage: 5.3 MB
```

Model Building

```
In [23]: # Dividing our data into input and output variables
X=pd.concat([data.iloc[:,1:]],axis=1)
y=data.iloc[:,0:1]

In [28]: # Logistic regression model
classifier=LogisticRegression()
classifier.fit(X,y)
```

Out[28]: LogisticRegression()


Model prediction and evaluation

```
In [25]: # Predict for x dataset
y_pred=classifier.predict(X)
y_pred

Out[25]: array([0, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [34]: ► confusion_matrix = confusion_matrix(y,y_pred)
confusion_matrix
fig, ax = plt.subplots(figsize=(8, 8))
sns.heatmap(confusion_matrix, cmap='magma', annot=True, fmt=".2f")
```

Out[34]: <AxesSubplot:>



```
In [35]: ► confusion_matrix
```

Out[35]: array([[39109, 813],
 [4015, 1274]], dtype=int64)

```
In [36]: ► # The model accuracy is calculated by (a+d)/(a+b+c+d)
(39109+1274)/(39109+813+4015+1274)
```

Out[36]: 0.8932118289796731

Model accuracy is : 89.31%


```
In [37]: ► classifier.predict_proba(X)[: ,1]
```

Out[37]: array([0.04534893, 0.01988662, 0.0116982 , ..., 0.83260083, 0.07658997,
 0.11785517])

```
In [40]: ▶ # ROC Curve plotting and finding AUC value
fpr, tpr, thresholds=roc_curve(y, classifier.predict_proba(X)[: ,1])
plt.plot(fpr, tpr, color='blue')
auc=roc_auc_score(y, y_pred)

plt.plot(fpr, tpr, color='blue', label='logit model(area = %0.2f)'%auc)
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
plt.ylabel('True Positive Rate')
plt.show()

print('auc accuracy:', auc)
```



auc accuracy: 0.6102562906535205

```
In [ ]: ▶
```